

biphy 1.0

Walker Pett

November 6, 2014

1 Introduction

1.1 Markovian Substitution Process

The process of gain and loss of binary characters is modeled as a continuous time Poisson process along each branch, where the rates of gain and loss on branch j of duration t_j are indicated by λ_j and μ_j respectively. The stochastic kernel of this substitution process on branch j , denoted \mathbf{Q}_j , can be expressed as follows:

$$\mathbf{Q}_j = \begin{bmatrix} -\lambda_j & \lambda_j \\ \mu_j & -\mu_j \end{bmatrix}$$

When not constrained to ultrametric trees, it is more natural to use the normalized rates, $\rho_j = \lambda_j + \mu_j$ and $\pi_j = \frac{\lambda_j}{\lambda_j + \mu_j}$, such that

$$\mathbf{Q}_j / \rho_j = \begin{bmatrix} -\pi_j & \pi_j \\ (1 - \pi_j) & -(1 - \pi_j) \end{bmatrix}$$

The parameters π_j and ρ_j can be interpreted as the stationary gene frequency and total rate of character replacement on branch j , respectively. Using this reparameterization, the branch lengths can then be re-expressed in units of the expected number of substitutions per site, $\nu_j = \rho_j t_j$

Note that because, in general, the entire substitution process along τ is not time-reversible, the placement of the root is not trivial. Furthermore, because the process is not stationary, we specify an independent parameter ϕ representing the stationary frequency in the ancestral sequence at the root of τ .

2 Basic usage

2.1 Starting a run

a run is initiated by providing the `biphy` command with a data file and a run name

```
biphy [options...] -d <data file> <run name>
```

2.2 Summarizing the output

`biphy` produces a trace file (`<run name>.trace`) and a treelist file (`<run name>.treelist`) suitable as inputs respectively to the `tracecomp` and `bpcomp` programs in the PhyloBayes package [3].

2.3 Restarting a run

In the event of a power failure or other error, a **biphy** run can be restarted using the run name without any other options:

```
biphy <run name>
```

2.4 Stream-reading mode

Stream-reading mode is activated when some options are specified, but a data file is not specified via the **-d** option. this mode is used only for performing posterior predictive simulations and computing cross-validation scores via the **-ppred** and **-cv** options:

```
biphy [-ppred -cv <test file>] <run name>
```

note that a run for which the **-s** option was not specified produces a stream that contains only one sample (the most recent sample). thus, activating stream-reading mode on such a run will give an error.

3 Detailed Options

General options

-d <datafile>

option for specifying the multiple sequence alignment to be analysed (see: Input Data Format section)

-x <every> [<until>]

specifies the saving frequency and (optional) the number of points after which the chain should stop. If this number is not specified, the chain runs forever. By definition, **-x 1** corresponds to the default saving frequency.

-t <tree file>

constrains the analysis to be conducted on the specified fixed tree

-o <outgroup file>

constrains trees to be rooted using the specified outgroup

-f

forces the program to overwrite an already existing run with same name.

Model options

-dollo

the stochastic Dollo model of Nicholls and Gray [4]. in this case, $\pi = 0$ and $\phi = 1$, and the likelihood is computed by summing over all possible “root” nodes for each character (the nodes where that character could have been born). under this model, branch lengths are expressed in expected numbers of losses per site. this option automatically activates the **-u 1** option

-h

time-homogeneous model. a single stationary frequency ϕ is assumed across the entire tree.

-nh

time-heterogeneous, hierarchical beta model. branch-wise stationary frequencies are drawn iid from a beta distribution with hyperparameters β_1 and β_2

-m <int>

time-heterogeneous, beta mixture model. branch-wise stationary frequencies are drawn iid from a mixture of <int> Beta distributions with hyperparameters β_1 and β_2

-dpp

time-heterogeneous, infinite mixture model. branch-wise stationary frequencies are drawn from a Dirichlet process with Exponential concentration parameter of mean 1, and with a Beta base distribution with hyperparameters β_1 and β_2

Additional priors

-dgam <int>

this option specifies that site-specific rates should be drawn from a $\text{Gamma}(\alpha, \alpha)$ distribution, discretized on <int> quantiles. the options **-dgam 0** or **-dgam 1** specify a constant rates model. the default is 4 rate categories.

-lexp

hierarchical exponential prior on branch lengths. branch lengths are drawn from an exponential distribution of mean μ , which is itself drawn from an exponential distribution of mean 1/10 .

-ldir

compound Dirichlet prior on branch lengths [5]. branch lengths are drawn from a Dirichlet distribution of concentration parameter 1, scaled by the tree length which is drawn from an exponential distribution of mean 1.

-rp <min> <max>

the root frequency ϕ is truncated on the (unnormalized) interval (**min,max**). applies only to the time homogeneous model, and the hierarchical beta model.

-rigid

the root frequency ϕ and the stationary frequencies on the branches subtending the root node of the tree are constrained to be equal. this can improve mixing due to the non-identifiability of these parameters in the unconstrained model [1]. applies only to time-heterogeneous models.

Correction for unobservable site-patterns

`-u <int>`

use this option to specify site-patterns that cannot be observed or which have been omitted from the dataset. this results in a correction to the likelihood that accounts for the absence of these site-patterns [2].

the option `<int>` is one of:

- 0 = no site-patterns have been omitted (default)
- 1 = constant absence site-patterns have been omitted
- 2 = constant presence site-patterns have been omitted
- 4 = singleton gains have been omitted
- 8 = singleton losses have been omitted

combinations of the above site-patterns are achieved by adding the indicated values. for example, `-u 3` indicates that constant sites have been omitted, and `-u 15` indicates that parsimony uninformative sites have been omitted.

MCMCMC options

`-n <int>`

this option specifies the number of Metropolis-coupled chains to use, with one chain per available thread. the default is 1 chain.

`-delta <float>`

`-sigma <float>`

these options specify the δ and σ tuning parameters for the calculation of chain heats during a parallel Metropolis-coupled run. the heat of chain k is computed as follows

$$\beta_k = (1 + \delta)^{-\sigma^k}$$

the default values are $\delta = 0.05$ and $\sigma = 1$. these values should be appropriate for small values of n (less than 10), but should be adjusted on a case by case basis to improve mixing, usually by increasing σ and/or decreasing δ as the number of chains is increased.

`-si <int>`

this option specifies the number of generations between proposals to swap Metropolis-coupled chains. the default value is 1.

Output options

`-e`

in addition to the Newick treelist output, a Nexus-formatted tree file with extra branch parameters is produced. This is useful for mapping branch-wise stationary frequencies to branches in tree visualisation software.

`-s`

the entire output of all chains is saved to the file `<run name>.stream`. This option is required if subsequent posterior predictive analysis or cross-validation checks are desired. note that the saving frequency of this stream is the single-chain saving frequency specified by the `-x` option, multiplied by the swap interval of the Metropolis-coupled chain specified by the `-si` option.

Model-checking options

the following options are to be used in stream-reading mode, which is activated by running `biphy` without specifying a data file via the `-d` option. in stream-reading mode, the parameter configuration of each sample is read sequentially from the stream of the indicated run, and computations are performed by conditioning on this parameter configuration

`-ppred`

using this option, an alignment is simulated at each sample point, then the frequency of “1” states is calculated for each species and saved to the file `<run name>.ppred`

`-cv <test file>`

using this option, the likelihood of the specified test alignment is calculated at each sample point, and saved to the file `<run name>.cv`

References

- [1] Joseph T Chang. Full Reconstruction of Markov Models on Evolutionary Trees : Identifiability and Consistency. *Mathematical Biosciences*, 13:51–73, 1996.
- [2] Joseph Felsenstein. Phylogenies from restriction sites: A maximum-likelihood approach. *Evolution*, 46(1):159–173, 1992.
- [3] Nicolas Lartillot and Hervé Philippe. A Bayesian mixture model for across-site heterogeneities in the amino-acid replacement process. *Molecular biology and evolution*, 21(6):1095–109, June 2004.
- [4] Geoff K. Nicholls and Russell D. Gray. Dated ancestral trees from binary trait data and their application to the diversification of languages. *Journal of the Royal Statistical Society: Series B (Statistical Methodology)*, 70(3):545–566, July 2008.
- [5] Bruce Rannala, Tianqi Zhu, and Ziheng Yang. Tail paradox, partial identifiability, and influential priors in Bayesian branch length inference. *Molecular biology and evolution*, 29(1):325–35, January 2012.